

Sparse Matrix Multiplication

Alok Ranjan

241010011

Department of Aerospace Engineering
IIT Kanpur

April 21, 2025

Serial: Core Mathematics

For matrices A and B , compute:

$$C_{ij} = \sum_{k=1}^{80} A_{ik} \times B_{kj}$$

Basic Optimization

Skip when $A_{ik} = 0$:

$$C_{ij} = \sum_{\substack{k=1 \\ A_{ik} \neq 0}}^{80} A_{ik} \times B_{kj}$$

- ▶ **Operations Saved:** 70% (from 512,000 to 153,600)
- ▶ **Key Idea:** Check before multiplying

Serial Code Explained

```
1  for (int i = 0; i < 80; i++) {           // For each
    output row
2    for (int k = 0; k < 80; k++) {         // For each
        element
3        if (A[i][k] != 0) {               // zero-check
4            for (int j = 0; j < 80; j++) {
5                C[i][j] += A[i][k] * B[k][j];
6            }
        }
    }
```

- ▶ **Loop Order:** $i \rightarrow k \rightarrow j$ optimal for row-major memory
- ▶ **Why This Works:** Skips entire inner loops when $A_{ik} = 0$

OpenMP: Parallel Mathematics

Extend serial approach by distributing rows:

$$C_{\text{thread}} = \sum_{\text{assigned rows}} A_{\text{thread}} \times B$$

Implementation

- ▶ Divide A 's rows equally among threads
- ▶ Each thread computes its portion independently

OpenMP Code

```
1 #pragma omp parallel for // parallel directive
2 for (int i = 0; i < 80; i++) {
3     // Same serial logic:
4     if (A[i][k] != 0) { // zero-check
5         for (int j = 0; j < 80; j++) {
6             C[i][j] += A[i][k] * B[k][j];
7         }
8     }
9 }
```

► How Distribution Works:

- Thread 0: rows 0-19
- Thread 1: rows 20-39
- etc.

► Why Safe: Threads write to different $C[i]$ rows

MPI: Data Distribution Mechanics

Scattering Matrix A:

- ▶ `MPI_Scatter` splits A by rows
- ▶ Example (4 processes):
 - ▶ Process 0: rows 0-19
 - ▶ Process 1: rows 20-39
 - ▶ Process 2: rows 40-59
 - ▶ Process 3: rows 60-79
- ▶ Each gets `localA[20][80]`

Broadcasting Matrix B:

- ▶ `MPI_Bcast` sends full B to all
- ▶ Each process gets `B[80][80]`
- ▶ Why? All need full B for multiplication

MPI: Computation & Gathering

Local Computation:

- ▶ Each process calculates:

$$\text{localC} = \text{localA} \times B$$

- ▶ Uses same zero-skipping as serial
- ▶ Works on its 20-row chunk

Gathering Results:

- ▶ `MPI_Gather` collects `localC` chunks
- ▶ Reassembles full C at Rank 0
- ▶ Like puzzle pieces coming together

MPI Code

```
1 // Rank 0 divides A
2 MPI_Scatter(A, 20*80, MPI_INT, localA, 20*80, MPI_INT, 0,
   MPI_COMM_WORLD);
3
4 // All get full B
5 MPI_Bcast(B, 80*80, MPI_INT, 0, MPI_COMM_WORLD);
6
7 // local computation
8 for (i = 0; i < 20; i++) {
9     if (localA[i][k] != 0) { // zero-check
10         for (j = 0; j < 80; j++) {
11             localC[i][j] += localA[i][k] * B[k][j];
12         }
13     }
14
15 // Rank 0 collects results
16 MPI_Gather(localC, 20*80, MPI_INT, C, 20*80, MPI_INT, 0,
   MPI_COMM_WORLD);
```


Execution Time Comparison

Implementation	2 Threads	4 Threads	8 Threads
Serial		0.0066 s	
OpenMP	0.002713 s	0.001528 s	0.004866 s
MPI	0.002752 s	0.002797 s	0.000853 s

Table: Execution times for different parallel implementations

- ▶ MPI with 8 threads showed the best performance (0.000853 s)
- ▶ OpenMP performance degraded with 8 threads (0.004866 s)
- ▶ Serial implementation was the slowest (0.0066 s)

Parallel and Serial $Ax=b$ linear Solvers

Implementing Jacobi, Gradient Descent and PETSc

Jacobi Method - Algorithm

Mathematical Formulation

Update each variable using previous iteration values:

$$x_i^{(k+1)} = \frac{1}{a_{ii}} \left(b_i - \sum_{j \neq i} a_{ij} x_j^{(k)} \right)$$

Algorithm 1 Jacobi Iteration

- 1: Initialize $x^{(0)} = 0$
- 2: **for** $k = 1$ to max_iterations **do**
- 3: **for** $i = 1$ to n **do**
- 4: $\sigma \leftarrow \sum_{j \neq i} a_{ij} x_j^{(k)}$
- 5: $x_i^{(k+1)} \leftarrow (b_i - \sigma) / a_{ii}$
- 6: **end for**
- 7: **if** $\|x^{(k+1)} - x^{(k)}\| < \epsilon$ **then break**
- 8:

Jacobi - Serial Code

```
1 // Main Jacobi loop
2 for (k = 0; k < MAX_ITER; k++) {
3     for (i = 0; i < N; i++) {
4         sum = 0;
5         for (j = 0; j < N; j++) {
6             if (j != i) sum += A[i][j] * x[j];
7         }
8         x_new[i] = (b[i] - sum) / A[i][i];
9     }
10    error = compute_error(x_new, x);
11    if (error < TOL) break;
12    // Update for next iteration
13    for (i = 0; i < N; i++) x[i] = x_new[i];
14 }
```

- ▶ Time: 0.000094 sec (16×16 matrix)
- ▶ Pros: Simple, easy to parallelize

Gradient Descent - Algorithm

Key Steps

1. Compute residual: $r = b - Ax$
2. Calculate step size: $\alpha = \frac{r^T r}{r^T A r}$
3. Update solution: $x = x + \alpha r$

Algorithm 2 Gradient Descent

- 1: $x \leftarrow 0$
 - 2: **repeat**
 - 3: $r \leftarrow b - Ax$
 - 4: $\alpha \leftarrow (r^T r)/(r^T A r)$
 - 5: $x \leftarrow x + \alpha r$
 - 6: **until** $\|r\| < \epsilon = 0$
-

Gradient Descent - Serial Code

```
1  for (iter = 0; iter < MAX_ITER; iter++) {
2      // Compute residual  $r = b - A*x$ 
3      mat_vec_mult(A, x, Ax);
4      vector_sub(b, Ax, r);
5
6      // Compute step size
7      dot_rr = vector_dot(r, r);
8      mat_vec_mult(A, r, Ar);
9      alpha = dot_rr / vector_dot(r, Ar);
10
11     // Update solution
12     vector_add_scaled(x, alpha, r, x);
13
14     if (sqrt(dot_rr) < TOL) break;
15 }
```

- ▶ Time: 0.000604 sec (16×16 matrix)
- ▶ Pros: Faster convergence for large systems

PETSc - Key Components

Matrix Setup

- ▶ `MatCreate()`: Creates matrix
- ▶ `MATAIJ`: Sparse format
- ▶ `Preallocation`: Critical for performance

Solver Setup

- ▶ `KSPCreate()`: Solver context
- ▶ `KSPCG`: Conjugate Gradient
- ▶ `KSPSetTolerances()`: Convergence

Why These Choices?

- ▶ Efficient for sparse systems
- ▶ Robust convergence
- ▶ Parallel-ready design

PETSc - Code Implementation

```
1 // 1. Matrix Setup
2 MatCreate(PETSC_COMM_WORLD, &A);
3 MatSetSizes(A, PETSC_DECIDE, PETSC_DECIDE, N, N);
4 MatSetType(A, MATAIJ); // Sparse format
5 MatMPIAIJSetPreallocation(A, 5, NULL, 5, NULL);
6
7 // 2. Vector Setup
8 VecCreate(PETSC_COMM_WORLD, &x);
9 VecSetSizes(x, PETSC_DECIDE, N);
10 VecDuplicate(x, &b);
11
12 // 3. Solver Configuration
13 KSPCreate(PETSC_COMM_WORLD, &ksp);
14 KSPSetType(ksp, KSPCG); // Conjugate Gradient
15 KSPSetTolerances(ksp, 1e-6, PETSC_DEFAULT, PETSC_DEFAULT,
16                  1000);
17
18 // 4. Solve System
19 KSPSetOperators(ksp, A, A);
20 KSPSolve(ksp, b, x);
```


PETSc - Behind the Scenes

- ▶ **Matrix Assembly:**
 - ▶ PETSc distributes matrix across MPI processes
 - ▶ Each process owns portion of rows
- ▶ **Solver Workflow:**
 - ▶ Builds Krylov subspace (CG: conjugate directions)
 - ▶ Applies preconditioner (default: block Jacobi)
 - ▶ Checks convergence criteria
- ▶ **Parallel Communication:**
 - ▶ MPI handles inter-process data exchange
 - ▶ Overlaps computation/communication

Execution Results

Implementation	Time (s)
Serial Jacobi	0.000094
Serial Gradient	0.000604
PETSc Jacobi	—
PETSc Gradient	—

- ▶ Jacobi: Fastest for small matrices
- ▶ Gradient: Better theoretical scaling
- ▶ PETSc: Implementation in progress

Thank You!